

802

automated checks, all passing, in the 11 suites below

31

pages checked by axe-core in every role: 0 violations

25/25

deployment checks against production (read-only; 3 upload checks skipped)

17 / 24

tables / foreign keys, asserted by the audit so the ERD cannot drift

| Suite                                                   | What it proves                                                                                                                                                                                         | Result  |
|---------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------|
| <b>Security and functional audit</b><br>npm run audit   | input validation, SQL injection, authentication, authorization, XSS, uploads, location privacy, report access and editing, matching lifecycle, the table and key counts                                | 382/382 |
| <b>Account lifecycle</b><br>scripts/auth_lifecycle.py   | registration, email verification, password reset and the sessions it ends, email change, rate limits, the password and name rules                                                                      | 106/106 |
| <b>Multi-device sessions</b><br>npm run multi-device    | independent sessions: suspension, lock and unlock take effect everywhere at once; a coordinator or administrator keeps one session at a time                                                           | 73/73   |
| <b>Sign-out privacy</b><br>npm run test:signout         | in a real Chrome: signing out and switching accounts leaves nothing of the last person on screen                                                                                                       | 24/24   |
| <b>Interface regressions</b><br>npm run test:ui         | in a real Chrome: moderation, registration feedback, phone menus, report actions, the edit freeze, historical labels, password requirements and strength, reopening, decisions asking first            | 66/66   |
| <b>Form feedback</b><br>npm run test:feedback           | in a real Chrome: Confirm password appears when the password qualifies and must be retyped; a submit from the bottom of a page ends on its answer; the date filter and controls at touch-laptop widths | 49/49   |
| <b>Contract tests</b><br>npm run test:contract          | what the form sends, what the API returns and what MySQL stores agree; dates in Philippine time; the password rule matches the one on the server                                                       | 36/36   |
| <b>Name and password rules</b><br>npm run test:identity | first and last name on the server; a password may not contain either; the same cases the browser copy is held to                                                                                       | 41/41   |
| <b>Place comparison</b><br>npm run test:city            | how two places compare when a report has no map pin                                                                                                                                                    | 11/11   |
| <b>Calendar</b><br>npm run test:calendar                | "today" is the Philippine date; the database speaks UTC                                                                                                                                                | 8/8     |
| <b>Matching log</b><br>npm run test:matching-log        | what matching writes to the log, with debugging off and on                                                                                                                                             | 6/6     |

WHY MORE THAN ONE KIND OF TEST

FINDING 1 · ENDPOINT TESTING

**A router fault no unit could show**  
Every handler in api/ was correct on its own, but api/index.php passed only the first two path segments to most of them, so a third was silently dropped: GET /api/matches/1/claims answered with the match. Only asking the running API for endpoints it does not have could find it. Fixed; audit cases EH-07 to EH-10.

FINDING 2 · MEASURING, NOT ONLY SCANNING

**Contrast that axe-core could not see**  
axe reported zero violations, but it cannot judge text over a photograph. Measuring the rendered pixels behind each line of text, at four screen widths, found three paragraphs under WCAG AA, the worst at 1.17:1, on pages axe had called clean. Fixed, then re-measured: every line now passes at every width.

EVERY FIX IS PROVEN RED FIRST

A regression test is run against the old code before the fix, and must fail there. Examples from the final fixes: the edit freeze (EM-03, 04, 05, 05b, 13 failed before), and refusing a reset to the current password (RS-1a–g, RS-4a).

WHERE TO FIND IT

docs/TESTING.md lists every suite, what it needs and what it changes. The numbers on this page are read from its section 3, verified 2 October 2026.

**Also:** lint clean · production build green · Docker image builds clean · schema.sql imports clean on MySQL 8.0 and 9.4 and matches the migrated database table for table. The suites run against the development copy: the audit and account tests reseed and change data, so production is checked read-only by the deployment verifier and by a manual walk-through.